


**[Network Layer] AsyncLLMClient vs vLLM Reference (N=128, R=32, Qwen2.5-7B on Clariden GH200) · N = 3 pairs, mean ± stdev**

**Throughput on vLLM workload**



**Per-request latency distribution**

