


[Network Layer] AsyncLLMClient vs vLLM Performance Reference (N=128, R=32, Qwen2.5-7B on DGX Spark) · N = 3 pairs, mean $\pm$ stdev

Throughput on vLLM workload



Per-request latency distribution

